

```

bigdata@bigdata-VirtualBox:~/A1$ export
HADOOP_CLASSPATH=$(HADOOP_HOME/bin/hadoop classpath)
bigdata@bigdata-VirtualBox:~/A1$ javac -classpath ${HADOOP_CLASSPATH}
solution3.java
bigdata@bigdata-VirtualBox:~/A1$ jar cf solution3.jar solution3*.class
bigdata@bigdata-VirtualBox:~/A1$ $HADOOP_HOME/bin/hadoop jar
solution3.jar solution3 sales.txt /user/bigdata/output
21/10/13 10:18:02 INFO client.RMPProxy: Connecting to ResourceManager at
bigdata-VirtualBox/10.0.2.15:8032
21/10/13 10:18:03 WARN mapreduce.JobResourceUploader: Hadoop command-line
option parsing not performed. Implement the Tool interface and execute
your application with ToolRunner to remedy this.
21/10/13 10:18:03 INFO input.FileInputFormat: Total input paths to
process : 1
21/10/13 10:18:03 INFO mapreduce.JobSubmitter: number of splits:1
21/10/13 10:18:03 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1634070937113_0022
21/10/13 10:18:03 INFO impl.YarnClientImpl: Submitted application
application_1634070937113_0022
21/10/13 10:18:03 INFO mapreduce.Job: The url to track the job:
http://bigdata-VirtualBox:8088/proxy/application_1634070937113_0022/
21/10/13 10:18:03 INFO mapreduce.Job: Running job: job_1634070937113_0022
21/10/13 10:18:10 INFO mapreduce.Job: Job job_1634070937113_0022 running
in uber mode : false
21/10/13 10:18:10 INFO mapreduce.Job:  map 0% reduce 0%
21/10/13 10:18:15 INFO mapreduce.Job:  map 100% reduce 0%
21/10/13 10:18:20 INFO mapreduce.Job:  map 100% reduce 100%
21/10/13 10:18:20 INFO mapreduce.Job: Job job_1634070937113_0022
completed successfully
21/10/13 10:18:20 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=88
        FILE: Number of bytes written=237407
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=164
        HDFS: Number of bytes written=95
        HDFS: Number of read operations=6
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=1
        Launched reduce tasks=1
        Data-local map tasks=1
        Total time spent by all maps in occupied slots (ms)=2614
        Total time spent by all reduces in occupied slots (ms)=2953
        Total time spent by all map tasks (ms)=2614
        Total time spent by all reduce tasks (ms)=2953
        Total vcore-milliseconds taken by all map tasks=2614
        Total vcore-milliseconds taken by all reduce tasks=2953
        Total megabyte-milliseconds taken by all map tasks=2676736
        Total megabyte-milliseconds taken by all reduce tasks=3023872
    Map-Reduce Framework
        Map input records=7
        Map output records=7
        Map output bytes=68
        Map output materialized bytes=88

```

```

Input split bytes=109
Combine input records=0
Combine output records=0
Reduce input groups=3
Reduce shuffle bytes=88
Reduce input records=7
Reduce output records=12
Spilled Records=14
Shuffled Maps=1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=124
CPU time spent (ms)=1290
Physical memory (bytes) snapshot=447455232
Virtual memory (bytes) snapshot=3828072448
Total committed heap usage (bytes)=319815680

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

```

```

import java.io.IOException;
import java.util.StringTokenizer;

```

```

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.FloatWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

```

```

public class solution3 {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "solution3");
        job.setJarByClass(solution3.class);
        job.setMapperClass(TokenizerMapper.class);
        job.setReducerClass(IntMinMaxReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable counter = new IntWritable(0);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                counter.set( Integer.parseInt(itr.nextToken()) );
                context.write(word, counter);
            }
        }
    }
}

```

```

public static class IntMinMaxReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result_min_max = new IntWritable();
    private IntWritable result_sum = new IntWritable();
    private IntWritable result_avg = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {
        int max = Integer.MIN_VALUE;
        int min = Integer.MAX_VALUE;
        int value;
        int total = 0;
        int counter = 0;

        for (IntWritable val : values) {
            value = val.get();
            total = total + value;

            counter++;

            if ( value > max )
                max = value;
            if ( value < min )
                min = value;
        }
        result_sum.set(total);
        context.write(key, result_sum);

        result_avg.set(total/counter);
        context.write(key, result_avg);

        result_min_max.set(max);
        context.write(key, result_min_max);
        result_min_max.set(min);
        context.write(key, result_min_max);
    }
}

```