

ISIT312 Big Data Management

# Physical Data Warehouse Design

Dr Fenghui Ren

School of Computing and Information Technology -  
University of Wollongong

# Physical Data Warehouse Design

## Outline

Techniques for Physical Data Warehouse Design

Materialized View

Indexes for Data Warehouses

Evaluation of Star Queries

Data Warehouse Partitioning

# Techniques for Physical Data Warehouse Design

## Materialized Views

- A view physically stored in the DB
- Typical problems: view update, view selection

## Indexing

- Used in Data Warehouse together with materialized views
- Specific for Data Warehouse: bitmap and join indexes

## Partitioning

- Divides the contents of a relational table into several files
- Horizontal and vertical partitioning

# Physical Data Warehouse Design

## Outline

Techniques for Physical Data Warehouse Design

Materialized View

Indexes for Data Warehouses

Evaluation of Star Queries

Data Warehouse Partitioning

# Materialized Views

**Materialized view** is a relational table that contains the rows that would be returned by the view definition - usually **SELECT** statement of SQL

If we consider relational views as **stored queries** then **materialized views** can be considered as **stored results**

**Materialized views** are created and used to reduce an amount of time needed to compute **SELECT** statements, for example join materialized views eliminate the needs to join the relational table

There are two ways how materialized view can be used:

- brute force method
- transparent query rewrite

**In brute force method** SQL is written to explicitly access the view

**Transparent query rewrite** method is applied when a query optimizer detects that a query can be computed against a materialized view instead of the source relational tables

# Materialized Views

**View maintenance** means that when the base relational tables are updated then a materialized view must be updated too

**Incremental view maintenance** means that updated view is computed from the individual modifications to the relational tables and not from the entire relational tables

Creating **materialized view**

Creating materialized view

```
CREATE MATERIALIZED VIEW MV_ORDERS
REFRESH ON COMMIT
ENABLE QUERY REWRITE
AS( SELECT O_ORDERKEY, O_CUSTKEY, O_TOTALPRICE, O_ORDERDATE
    FROM ORDERS
    WHERE O_ORDERDATE > TO_DATE( '31-DEC-1986', 'DD-MON-YYYY' ) );
```

Direct access to **materialized view**

Directaccess to materialized view

```
SELECT *
FROM MV_ORDERS
WHERE O_ORDERDATE = TO_DATE( '01-JAN-1992', 'DD-MON-YYYY' )
```

# Materialized Views

Access to **materialized view** through query rewriting

Indirect access to materialized view through query rewriting

```
SELECT O_ORDERKEY, O_CUSTKEY, O_TOTALPRICE, O_ORDERDATE
FROM ORDERS
WHERE O_ORDERDATE > TO_DATE('31-DEC-1986','DD-MON-YYYY');
```

- The results from **EXPLAIN PLAN** statement

Query processing plan

PLAN\_TABLE\_OUTPUT

|     |                              |           |      |       |         |          |
|-----|------------------------------|-----------|------|-------|---------|----------|
| 0   | SELECT STATEMENT             |           | 108K | 2539K | 507 (1) | 00:00:01 |
| * 1 | MAT_VIEW REWRITE ACCESS FULL | MV_ORDERS | 108K | 2539K | 507 (1) | 00:00:01 |

Predicate Information (identified by operation id):

```
1 - filter("MV_ORDERS"."O_ORDERDATE">TO_DATE(' 1986-12-31 00:00:00',
        'syyyymm-dd hh24:mi:ss'))
```

# Physical Data Warehouse Design

## Outline

Techniques for Physical Data Warehouse Design

Materialized View

Indexes for Data Warehouses

Evaluation of Star Queries

Data Warehouse Partitioning



# Indexes for Data Warehouses

An index provides a quick way to locate data of interest

Sample query

```
SELECT *  
FROM EMPLOYEE  
WHERE EmployeeKey = 007;
```

SELECT statement with equality condition in WHERE clause

With the help of an index over a column **EmployeeKey** (primary key in **EMPLOYEE** table), a single disk block access will suffice to answer the query

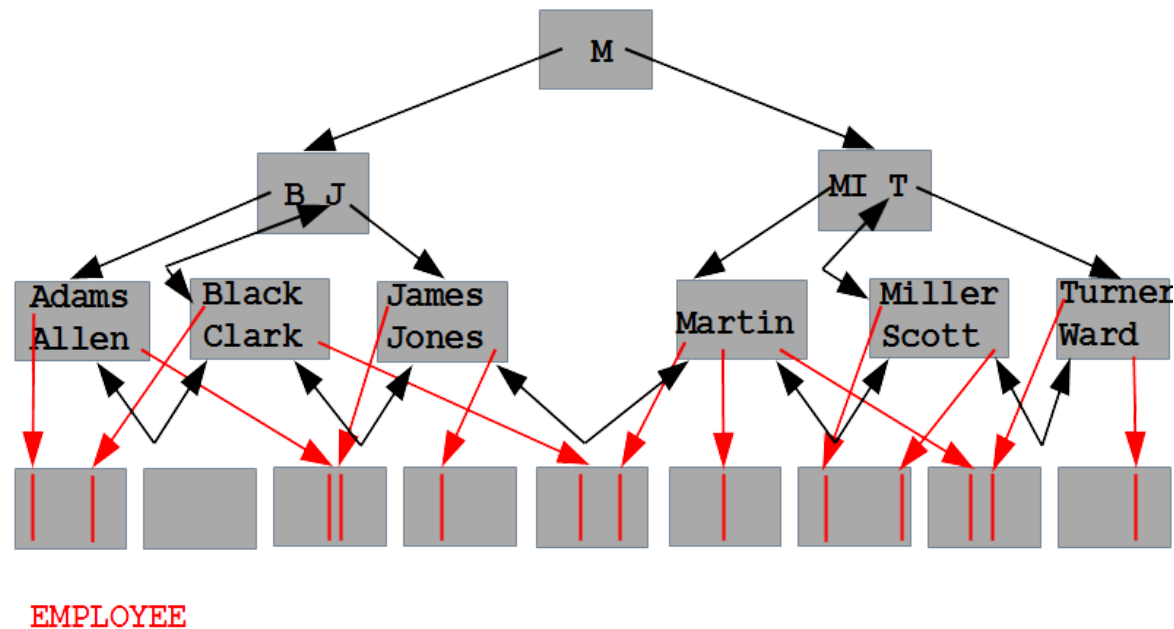
Without this index, we should perform a complete scan of table **EMPLOYEE**

**Drawback:** Almost every update on an indexed attribute also requires an index update

Too many indexes may degrade performance

Most popular indexing techniques in relational databases include **B\*-trees** and **bitmap indexes**

# B\*-tree index implementation



B\*-tree can be traversed either:

- vertically from root to leaf level of a tree
- horizontally either from left corner of leaf level to right corner of leaf level or the opposite
- vertically and later on horizontally either towards left lower corner or right lower corner of leaf level

# Bitmap Indexes

| ProductKey | ProductName | QuantityPerUnit | UnitPrice | Discontinued | CategoryKey |
|------------|-------------|-----------------|-----------|--------------|-------------|
| p1         | prod1       | 25              | 60        | No           | c1          |
| p2         | prod2       | 45              | 60        | Yes          | c1          |
| p3         | prod3       | 50              | 75        | No           | c2          |
| p4         | prod4       | 50              | 100       | Yes          | c2          |
| p5         | prod5       | 50              | 120       | No           | c3          |
| p6         | prod6       | 70              | 110       | Yes          | c4          |

**Product** dimension table

|    | 25 | 45 | 50 | 70 |
|----|----|----|----|----|
| p1 | 1  | 0  | 0  | 0  |
| p2 | 0  | 1  | 0  | 0  |
| p3 | 0  | 0  | 1  | 0  |
| p4 | 0  | 0  | 1  | 0  |
| p5 | 0  | 0  | 1  | 0  |
| p6 | 0  | 0  | 0  | 1  |

Bitmap index for attribute **QuantityPerUnit**

|    | 60 | 75 | 100 | 110 | 120 |
|----|----|----|-----|-----|-----|
| p1 | 1  | 0  | 0   | 0   | 0   |
| p2 | 1  | 0  | 0   | 0   | 0   |
| p3 | 0  | 1  | 0   | 0   | 0   |
| p4 | 0  | 0  | 1   | 0   | 0   |
| p5 | 0  | 0  | 0   | 0   | 1   |
| p6 | 0  | 0  | 0   | 1   | 0   |

Bitmap index for attribute **UnitPrice**

# Bitmap Indexes: Example

Products having **between 45 and 55** pieces per unit, **and** with a **unit price between 100 and 200**

|    | 45 | 50 | OR1 |
|----|----|----|-----|
| p1 | 0  | 0  | 0   |
| p2 | 1  | 0  | 1   |
| p3 | 0  | 1  | 1   |
| p4 | 0  | 1  | 1   |
| p5 | 0  | 1  | 1   |
| p6 | 0  | 0  | 0   |

OR for QuantityPerUnit

|    | 100 | 110 | 120 | OR2 |
|----|-----|-----|-----|-----|
| p1 | 0   | 0   | 0   | 0   |
| p2 | 0   | 0   | 0   | 0   |
| p3 | 0   | 0   | 0   | 0   |
| p4 | 1   | 0   | 0   | 1   |
| p5 | 0   | 0   | 1   | 1   |
| p6 | 0   | 1   | 0   | 1   |

OR for UnitPrice

|    | OR1 | OR2 | AND |
|----|-----|-----|-----|
| p1 | 0   | 0   | 0   |
| p2 | 1   | 0   | 0   |
| p3 | 1   | 0   | 0   |
| p4 | 1   | 1   | 1   |
| p5 | 1   | 1   | 1   |
| p6 | 0   | 1   | 0   |

AND operation

# Indexes for Data Warehouses: Requirements

## Symmetric partial match queries

- All dimensions of the cube should be symmetrically indexed, to be searched simultaneously

## Indexing at multiple levels of aggregation

- Summary tables must be indexed in the same way as base nonaggregated tables

## Efficient batch update

- The refreshing time of a data warehouse must be considered when designing the indexing schema

## Sparse data

- Typically, only 20% of the cells in a data cube are nonempty
- The indexing schema must deal efficiently with sparse and nonsparse data

# Physical Data Warehouse Design

## Outline

Techniques for Physical Data Warehouse Design

Materialized View

Indexes for Data Warehouses

Evaluation of Star Queries

Data Warehouse Partitioning

# Star Queries

Queries over star schemas are called **star queries**

Join the fact table with the dimension tables

A typical star query: total sales of discontinued products, by customer name and product name

Star query

```
SELECT ProductName, CustomerName, SUM(SalesAmount)
FROM Sales S, Customer C, Product P
WHERE S.CustomerKey = C.CustomerKey AND S.ProductKey = P.ProductKey AND
      P.Discontinued = 'Yes'
GROUP BY C.CustomerName, P.ProductName;
```

Three basic steps to evaluate the query:

- (1) Evaluation of the join conditions
- (2) Evaluation of the selection conditions over the dimensions
- (3) Aggregation of the tuples that passed the filter

# Evaluation of Star Queries with Bitmap Indexes: Example

| Product Key | Product Name | ... | Discontinued | ... |
|-------------|--------------|-----|--------------|-----|
| p1          | prod1        | ... | No           | ... |
| p2          | prod2        | ... | Yes          | ... |
| p3          | prod3        | ... | No           | ... |
| p4          | prod4        | ... | Yes          | ... |
| p5          | prod5        | ... | No           | ... |
| p6          | prod6        | ... | Yes          | ... |

Product table

| Yes | No |
|-----|----|
| 0   | 1  |
| 1   | 0  |
| 0   | 1  |
| 1   | 0  |
| 0   | 1  |
| 1   | 0  |

Bitmap for Discontinued

| Customer Key | Customer Name | Address          | Postal Code | ... |
|--------------|---------------|------------------|-------------|-----|
| c1           | cust1         | 35 Main St.      | 7373        | ... |
| c2           | cust2         | Av. Roosevelt 50 | 1050        | ... |
| c3           | cust3         | Av. Louise 233   | 1080        | ... |
| c4           | cust4         | Rue Gabrielle    | 1180        | ... |

Customer table

| Product Key | Customer Key | Time Key | Sales Amount |
|-------------|--------------|----------|--------------|
| p1          | c1           | t1       | 100          |
| p1          | c2           | t1       | 100          |
| p2          | c2           | t2       | 100          |
| p2          | c2           | t3       | 100          |
| p3          | c3           | t3       | 100          |
| p4          | c3           | t4       | 100          |
| p5          | c4           | t5       | 100          |

Sales fact table

| c1 | c2 | c3 | c4 |
|----|----|----|----|
| 1  | 0  | 0  | 0  |
| 0  | 1  | 0  | 0  |
| 0  | 1  | 0  | 0  |
| 0  | 1  | 0  | 0  |
| 0  | 0  | 1  | 0  |
| 0  | 0  | 1  | 0  |
| 0  | 0  | 1  | 0  |
| 0  | 0  | 0  | 1  |

Bitmap for CustomerKey

| p1 | p2 | p3 | p4 | p5 | p6 |
|----|----|----|----|----|----|
| 1  | 0  | 0  | 0  | 0  | 0  |
| 1  | 0  | 0  | 0  | 0  | 0  |
| 0  | 1  | 0  | 0  | 0  | 0  |
| 0  | 1  | 0  | 0  | 0  | 0  |
| 0  | 0  | 1  | 0  | 0  | 0  |
| 0  | 0  | 0  | 1  | 0  | 0  |
| 0  | 0  | 0  | 0  | 1  | 0  |
| 0  | 0  | 0  | 0  | 0  | 1  |

Bitmap for ProductKey

| Yes | No |
|-----|----|
| 0   | 1  |
| 0   | 1  |
| 1   | 0  |
| 1   | 0  |
| 0   | 1  |
| 1   | 0  |
| 0   | 1  |

Bitmap join index for Discontinued



# Evaluation of Star Queries using Bitmap Indexes

Evaluation of star query requires

- a B+ tree over **CustomerKey** and **ProductKey**
- Bitmap indexes on the foreign key columns in **Sales** and on **Discontinued** in **Product**

Example of query evaluation

- (1) Obtain the record numbers of the records that satisfy the condition **Discontinued = 'Yes'**
- Answer: Records with **ProductKey** values **p2** , **p4** , and **p6**
- (2) To access the bitmap vectors in **Sales** with these labels perform a join between **Product** and **Sales**
- (3) Vectors labeled **p2** and **p4** match, no fact record for **p6**
- (4) Obtain the values for the **CustomerKey** in these records (**c2** and **c3** )
- (5) Use B+-tree index on **ProductKey** and **CustomerKey** to find the names of products and customers
- (6) Answer: (**cust2,prod2,200**) and (**cust3,prod4,100**)

# Physical Data Warehouse Design

## Outline

Techniques for Physical Data Warehouse Design

Materialized View

Indexes for Data Warehouses

Evaluation of Star Queries

Data Warehouse Partitioning

# Data Warehouse Partitioning

**Partitioning** (or **fragmentation**) divides a table into smaller data sets (each one called a partition)

Applied to tables and indexes

Vendors provide several different partitioning methods

**Vertical partitioning** splits the attributes of a table into groups that can be independently stored

- E.g., most often used attributes are stored in one partition, less often used attributes in another one
- More records fit into main memory, reducing their processing time

**Horizontal partitioning** divides a table into smaller tables with same structure than the full table

- For example, if some queries require the most recent data, partition horizontally according to time

# Queries over Partitioned Databases

**Partition pruning** is the typical way of improving query performance using partitioning

Example: A **Sales** fact table in a warehouse can be partitioned by month

A query requesting orders for a single month only needs to access the partition of such a month

**Joins** also enhanced by using partitioning:

- When the two tables are partitioned on the join attributes
- When the reference table is partitioned on its primary key
- Large join is broken down into smaller joins

# Partitioning Strategies

Three partitioning strategies: Range partitioning, hash partitioning, and list partitioning

**Range partitioning** maps records to partitions based on ranges of values of the partitioning key

Time dimension is a natural candidate for range partitioning

Example: A table with a **date** column defined as the partitioning key

- **January-2012** partition will contain rows with key values from January 1 to January 31, 2012

**Hash partitioning** uses a hashing algorithm over the partitioning key to map records to partitions

- Hashing algorithm distributes rows among partitions in a uniform fashion, yielding, ideally, partitions of the same size
- Typically used when partitions are distributed in several devices, and when data are not partitioned based on time

# Partitioning Strategies

**List partitioning** specifies a list of values for the partitioning key

Some vendors (e.g. Oracle) support the notion of **composite partitioning**, combining the basic data distribution methods

Thus, a table can be range partitioned, and each partition can be subdivided using hash partitioning

# References

A. VAISMAN, E. ZIMANYI, Data Warehouse Systems: Design and Implementation, Chapter 7 Physical Data Warehouse Design, Springer Verlag, 2014